

# Inference for Data Visualization

Christof Seiler

Stanford University, Spring 2016, Stats 205

# Introduction



- ▶ Exploratory data analysis is usually not parametric
- ▶ For instance, in Principle Component Analysis (PCA), we do not assume any parametric model (the data doesn't need to be normally distributed)
- ▶ What is described by PCA is a decomposition of the data into Principle Components (PCs) along which the variance is maximized after projecting the data
- ▶ But, as we have seen in this course, it is in general not necessary to assume a parametric model for inference

# Introduction

- ▶ We successfully used ranks that allowed to remove the normality assumptions in one and two-sample tests
- ▶ We successfully used the bootstrap to sample from the empirical distribution and construct confidence intervals
- ▶ We successfully used permutation tests for hypothesis testing
- ▶ In all these examples we have found ways to make inference
- ▶ Is this possible for data visualization?
- ▶ That's the topic for today

# Magical Thinking


- ▶ Professional statisticians and other scientists with statistical training were asked “How associated the two variables were”



Source: Diaconis (1983)

- ▶ Most of the subjects judged left plot as more associated than right plot (the same data points)
- ▶ Rescaling can shift the perceived association by 10 to 15%

# Inference for Plots: The Lineup



Source: Buja et al. (2009)

## Inference for Plots: The Lineup

- ▶ Generate 19 null plots
- ▶ Arrange all 19 plots and insert the real data at random location
- ▶ Ask human viewer to single out the real plot
- ▶ Under the null hypothesis that all plots are the same, there is a one in 20 chance to single out the real one
- ▶ If the viewer chooses the plot of the real data, then the discovery can be assigned a  $p$ -value of  $1/20 = 0.05$
- ▶ Larger number of null plots could yield a smaller  $p$ -value
- ▶ But there is a limit of how many plots a human can consider

## Inference for Plots: The Lineup


- ▶ This protocol can be repeated with multiple independently recruited viewers
- ▶ Consider  $K$  viewers and  $k \leq K$  selected the plot of the real data
- ▶ Then the combined  $p$ -value is probability  $P(X \geq k)$  following a binomial distribution with  $K$  trials and success probability  $1/20$
- ▶ Can be as small as  $0.05^K$  if all viewers picked the plot of the real data


## Inference for Plots: The Lineup (Example)

- ▶ Example comes from Boyer & Savageau (1984) where cities across the USA were rated in 1984
- ▶ Question: Is 'Climate-Terrain' associated to 'Housing'?
- ▶ Low values on 'Climate-Terrain' imply uncomfortable temperatures (either hot or cold)
- ▶ High values of 'Housing' indicate a higher cost of owning a single family residence

## Inference for Plots: The Lineup (Example)

- ▶ The null hypothesis for this example is  
 $H_0$ : Housing is independent of Climate-Terrain
- ▶ The null plots are generated by permuting the values of the variable Housing
- ▶ Pick out the plot of the real data: Is any plot different from the others?
- ▶ Plots on next slide are taken from Buja et al. (2009)







- ▶ In class experiment



```
# number of students  
K = 8  
# number of correct picks  
k = 2  
pvalue = sum(dbinom(k:K,K,1/20)); pvalue
```

```
## [1] 0.05724465
```

## Inference for Plots: The Lineup (Example)

- ▶ HSBC (The Hongkong and Shanghai Banking Corporation) daily stock returns
  - ▶ two panels, the first showing the 2005 data only,
  - ▶ the second the more extensive 1998–2005 data
- ▶ In each panel, select which plot is the most different
- ▶ Plots on next slide are taken from Buja et al. (2009)





- ▶ In class experiment

```
# number of students  
K = 8  
# number of correct picks  
k = 4  
pvalue = sum(dbinom(k:K,K,1/8)); pvalue
```

```
## [1] 0.01124781
```

## Inference for Plots: The Lineup (Example)

- ▶ For 2005, the viewer should have had difficulty selecting the real data
  - ▶ This is a year of low and stable volatility
- ▶ For 1998–2005, it should be easy
  - ▶ features two volatility bursts
  - ▶ one in 1998 due to the Russian bond default and the LTCM collapse
  - ▶ the other in 2001 due to the 9/11 event
  - ▶ after, volatility stabilizes at a low level

# Principal Component Analysis

- ▶ Principal Component Analysis (PCA) is a data exploration tool
- ▶ PCA finds a low-dimensional subspace that minimizes the distances between projections points and subspace
- ▶ Consider observations  $x_1, x_2, \dots, x_n$
- ▶ Center and combine them in matrix  $X$  of dimension  $p \times n$
- ▶ PCA solves this minimization problem with  $\langle v_1, v_1 \rangle = 1$

$$\hat{v}_1 = \underset{v_1}{\text{maximize}} \left\{ \text{Var}(Xv_1) \right\}$$


- ▶ And for  $v_2$  with  $\langle v_1, v_2 \rangle = 0$  and  $\langle v_2, v_2 \rangle = 1$

$$\hat{v}_2 = \underset{v_2}{\text{maximize}} \left\{ \text{Var}(Xv_2) \right\}$$






- ▶ Keep going the same way until  $\hat{v}_1, \dots, \hat{v}_q$  have been collected and put them in  $\hat{V}_q$  of dimensions  $p \times q$

# Principal Component Analysis (Example)

Condylar process and head




Condylar process and head



Source: [www.aofoundation.org](http://www.aofoundation.org)

# Principal Component Analysis (Example)



Source: S., Pennec, and Reyes 2012

# Principal Component Analysis (Example)

Two animations of mandible “eigenanatomy”:

- ▶ <http://christofseiler.github.io/phd/>

# Bootstrap PCA




- ▶ Two ways to bootstrap PCA in case of random rows  $X$
- ▶ Partial bootstrap and total bootstrap
- ▶ Partial bootstrap:
  - ▶ Project  $B$  replications onto initial subspace
  - ▶ Initial subspace is obtained by PCA on original  $X$
  - ▶ Underestimates variation of parameters (Milan 1995)
- ▶ Total bootstrap:
  - ▶ Perform new PCA on each replication
  - ▶ Problem: Need to align PCA's
  - ▶ Nuisance variations: reflections and rotations

## Bootstrap PCA

- ▶ For the total bootstrap, need to align PCA's
- ▶ This is usually done using Procrustes analysis
- ▶ Procrustes refers to a bandit from Greek mythology who made his victims fit his bed by stretching their limbs (or cutting them off)
- ▶ Procrustes analysis is used in statistical shape analysis to compare aligned shapes after removing "nuisance" parameters:
  - ▶ translation in space
  - ▶ rotation in space
  - ▶ sometimes scaling of the objects


# Bootstrap PCA

- ▶ Shape example: landmarks for the human spine



## Bootstrap PCA

- ▶ Same idea can be applied to align projected observations
- ▶ In PCA, shapes are the projected observations onto the lower dimensional subspace spanned by say PC1 and PC2



Source: Josse, Wager, and Husson (2014)

## Bootstrap PCA

- ▶ Collecting  $B$  bootstrap sampled PCA's by resampling rows of data matrix  $X$

$$\hat{V}_q^{*1}, \dots, \hat{V}_q^{*B}$$

- ▶ Align all the projected point set using Procrustes alignment
- ▶ Meaning, we find rotation ( $R^T R = I$ )

$$\hat{R}^b = \underset{R}{\text{minimize}} = \left\{ \|X^{*1} \hat{V}_q^{*1} - X^{*b} \hat{V}_q^{*b} R\|^2 \right\}$$

- ▶ and apply rotation to projected data points

$$X^{*b} \hat{V}_q^{*b} \hat{R}^{*b}$$


- ▶ Overlay points and draw contours around it

# Parametric Bootstrap PCA

- ▶ In case of fixed rows and columns  $X$ , we can use parametric bootstrap
- ▶ It is good alternative when the model is too difficult or before the asymptotics regime
- ▶ Steps:
  1. Perform PCA on  $X$  to estimate  $\hat{V}_p$
  2. Estimate error  $\sigma^2$  from residual matrix  $\epsilon_{n \times p} = X - \hat{V}_q \hat{V}_q^T X$   
(assume elementwise iid normal noise)
  3. Bootstrap  $1, \dots, B$ :
    - ▶ Draw  $\epsilon_{ij}^{*b}$  from  $N(0, \hat{\sigma}^2)$
    - ▶ Generate new matrix  $X^{*b} = \hat{V}_q \hat{V}_q^T X + \epsilon^{*b}$
    - ▶ Perform PCA on  $X^{*b}$

# Parametric Bootstrap PCA (Example)


- ▶ Consumers describe 10 white wines with 15 sensory attributes
- ▶ Consumers score wines between 1 and 10 for each attribute
- ▶ Collect averages across consumers in  $10 \times 15$  matrix  $X$



Source: Josse et al.

# Parametric Bootstrap PCA (Example)

- With bootstraped confidence ellipses



Source: Josse et al.

## References

- ▶ Diaconis (1983). Theories of Data Analysis: From Magical Thinking Through Classical Statistics
- ▶ Buja, Cook, Hofmann, Lawrence, Lee, Swayne, and Wickham (2009). Statistical Inference for Exploratory Data Analysis and Model Diagnostics
- ▶ Milan and Whittaker (1995). Application of the Parametric Bootstrap to Models that Incorporate a Singular Value Decomposition
- ▶ Josse, Wager, and Husson (2014). Confidence Areas for Fixed-Effects PCA
- ▶ Seiler, Pennec, and Reyes (2012). Capturing the Multiscale Anatomical Shape Variability with Polyaffine Transformation Trees